

Module 3: R Manipulation

Instructor: Anjali Silva, PhD

TA: Tia Harrison, MSc

Data Sciences Institute, University of Toronto

2022-06-22

Course Documents

- Visit: <https://github.com/anjalisilva/IntroductionToR>
- All course material will be available via IntroductionToR GitHub repository (<https://github.com/anjalisilva/IntroductionToR>). Folder structure is as follows:
 - Lessons - All files: This folder contains all files.
 - **Lessons - Data only**: This folder contains data only.
 - **Lessons - Lesson Plans only**: This folder contains lesson plans only.
 - **Lessons - PDF only**: This folder contains slide PDFs only.
 - README - README file
 - .gitignore - Files to ignore specified by instructor

Course Contacts

- Instructor: Anjali Silva Email: a.silva@utoronto.ca (Must use the subject line DSI-IntroR. E.g., DSI-IntroR: Inquiry about Lecture I.)
- TA: Tia Harrison Email: tia.harrison@mail.utoronto.ca

Overview

- Filtering (Wickham and Grolemund, 2017 Chapter 5, Timbers et al. 2021, Chapter 3.6)
- Arranging (Wickham and Grolemund, 2017 Chapter 5)
- Selecting (Wickham and Grolemund, 2017 Chapter 5, Timbers et al. 2021, Chapter 3.5)
- The pipe (Wickham and Grolemund, 2017 Chapter 5 & 18; Timbers et al. 2021, Chapter 3.8)
- Mutating (Wickham and Grolemund, 2017 Chapter 5, Timbers et al. 2021, Chapter 3.7, 3.10)
- Summarising (Wickham and Grolemund, 2017 Chapter 5, Timbers et al. 2021, Chapter 3.9)
- Grouping (Wickham and Grolemund, 2017 Chapter 5)
- Cleaning (Alexander, 2022, Chapter 11)

Take a look

```
glimpse(ads_data)
```

```
## Rows: 1,460
## Columns: 52
## $ StartDate      <dtm> 2019-06-14 09:43:20, 2019-06-14 09:43:...
## $ EndDate        <dtm> 2019-06-14 09:44:30, 2019-06-14 09:44:...
## $ Status         <dbl+lbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,...
## $ Progress       <dbl> 100, 100, 100, 100, 100, 100, 100, 100, 100,...
## $ Duration__in_seconds_ <dbl> 70, 105, 88, 109, 109, 70, 99, 105, 124...
## $ Finished       <dbl+lbl> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,...
## $ RecordedDate   <dtm> 2019-06-14 09:44:31, 2019-06-14 09:44:...
## $ ResponseId     <chr> "R_11dq3s9btLX57LD", "R_DRWZdBOugPUKqGt...
## $ DistributionChannel <chr> "anonymous", "anonymous", "anonymous", ...
## $ UserLanguage   <chr> "EN", "EN", "EN", "EN", "EN", "EN", "EN...
## $ Consent        <dbl+lbl> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,...
## $ Pol_7          <dbl+lbl> 5, 3, 1, 2, 6, 4, 6, 4, 2, 5, 4, 1,...
## $ W2_Knowledge   <dbl+lbl> 2, 2, 4, 1, 3, 2, 3, 3, 3, 3, 3, 1,...
## $ Gender         <dbl+lbl> 2, 1, 2, 1, 1, 1, 1, 2, 1, 1, 1, 2,...
## $ Race           <dbl+lbl> 1, 1, 1, 1, 1, 3, 2, 1, 1, 1, 1, 2,...
## $ W1_Feeling_1   <dbl> 2, 1, 4, 3, 3, 3, 6, -6, 4, 1, 3, -1, 3...
## $ W1_Actions_1_1 <dbl+lbl> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA,...
## $ W1_Actions_1_2 <dbl+lbl> 1, NA, NA, 1, NA, NA, NA, NA, 1,...
## $ W1_Actions_1_3 <dbl+lbl> NA, NA, 1, NA, NA, 1, 1, NA, NA, 4/56
```

Filtering

Filtering allows us to select rows based on specific traits

```
filter(ads_data, Duration__in_seconds_ < 100)
```

```
## # A tibble: 41 × 52
```

```
##   StartDate      EndDate      Status Progress
##   <dtm>         <dtm>         <dbl+lbl>   <dbl>
```

```
## 1 2019-06-14 09:43:20 2019-06-14 09:44:30 0 [IP Address] 100
```

```
## 2 2019-06-14 09:43:29 2019-06-14 09:44:58 0 [IP Address] 100
```

```
## 3 2019-06-14 09:44:00 2019-06-14 09:45:11 0 [IP Address] 100
```

```
## 4 2019-06-14 09:43:32 2019-06-14 09:45:12 0 [IP Address] 100
```

```
## 5 2019-06-14 09:43:48 2019-06-14 09:45:25 0 [IP Address] 100
```

```
## 6 2019-06-14 09:44:24 2019-06-14 09:45:26 0 [IP Address] 100
```

```
## 7 2019-06-14 09:43:50 2019-06-14 09:45:29 0 [IP Address] 100
```

```
## 8 2019-06-14 09:44:15 2019-06-14 09:45:42 0 [IP Address] 100
```

```
## 9 2019-06-14 09:44:30 2019-06-14 09:45:58 0 [IP Address] 100
```

```
## 10 2019-06-14 09:44:36 2019-06-14 09:46:05 0 [IP Address] 100
```

```
## # ... with 31 more rows, and 48 more variables:
```

```
## #   Duration__in_seconds_ <dbl>, Finished <dbl+lbl>,
```

```
## #   RecordedDate <dtm>, ResponseId <chr>, DistributionChannel <chr>,
```

```
## #   UserLanguage <chr>, Consent <dbl+lbl>, Pol_7 <dbl+lbl>,
```

```
## #   W2_Knowledge <dbl+lbl>, Gender <dbl+lbl>, Race <dbl+lbl>,
```

```
## #   W1_Feeling_1 <dbl>, W1_Actions_1_1 <dbl+lbl>,
```

Arranging

Arranging allows us to sort the order of the table by a certain column

```
arrange(ads_data, Duration__in_seconds_)
```

```
## # A tibble: 1,460 × 52
##   StartDate      EndDate      Status Progress
##   <dtm>         <dtm>         <dbl+lbl>    <dbl>
## 1 2019-06-14 09:58:11 2019-06-14 09:59:01 0 [IP Address]    100
## 2 2019-06-14 09:44:24 2019-06-14 09:45:26 0 [IP Address]    100
## 3 2019-06-14 09:43:20 2019-06-14 09:44:30 0 [IP Address]    100
## 4 2019-06-14 09:44:00 2019-06-14 09:45:11 0 [IP Address]    100
## 5 2019-06-14 09:52:10 2019-06-14 09:53:26 0 [IP Address]    100
## 6 2019-06-14 09:45:57 2019-06-14 09:47:13 0 [IP Address]    100
## 7 2019-06-14 09:50:37 2019-06-14 09:51:53 0 [IP Address]    100
## 8 2019-06-14 09:45:49 2019-06-14 09:47:08 0 [IP Address]    100
## 9 2019-06-14 10:10:25 2019-06-14 10:11:45 0 [IP Address]    100
## 10 2019-06-14 09:53:33 2019-06-14 09:54:54 0 [IP Address]    100
## # ... with 1,450 more rows, and 48 more variables:
## #   Duration__in_seconds_ <dbl>, Finished <dbl+lbl>,
## #   RecordedDate <dtm>, ResponseId <chr>, DistributionChannel <chr>,
## #   UserLanguage <chr>, Consent <dbl+lbl>, Pol_7 <dbl+lbl>,
## #   W2_Knowledge <dbl+lbl>, Gender <dbl+lbl>, Race <dbl+lbl>,
## #   W1_Feeling_1 <dbl>, W1_Actions_1_1 <dbl+lbl>,
```

Selecting

Selecting allows us to pick certain columns

```
select(ads_data, RecordedDate)
```

```
## # A tibble: 1,460 × 1
##   RecordedDate
##   <dtm>
## 1 2019-06-14 09:44:31
## 2 2019-06-14 09:44:58
## 3 2019-06-14 09:44:59
## 4 2019-06-14 09:45:00
## 5 2019-06-14 09:45:01
## 6 2019-06-14 09:45:12
## 7 2019-06-14 09:45:12
## 8 2019-06-14 09:45:13
## 9 2019-06-14 09:45:13
## 10 2019-06-14 09:45:16
## # ... with 1,450 more rows
```

Selecting

We can also remove columns

```
select(ads_data, -Consent, -DistributionChannel)
```

```
## # A tibble: 1,460 × 50
```

```
##   StartDate      EndDate      Status Progress
##   <dtm>         <dtm>         <dbl+lbl>    <dbl>
## 1 2019-06-14 09:43:20 2019-06-14 09:44:30 0 [IP Address]    100
## 2 2019-06-14 09:43:11 2019-06-14 09:44:57 0 [IP Address]    100
## 3 2019-06-14 09:43:29 2019-06-14 09:44:58 0 [IP Address]    100
## 4 2019-06-14 09:43:10 2019-06-14 09:45:00 0 [IP Address]    100
## 5 2019-06-14 09:43:11 2019-06-14 09:45:00 0 [IP Address]    100
## 6 2019-06-14 09:44:00 2019-06-14 09:45:11 0 [IP Address]    100
## 7 2019-06-14 09:43:32 2019-06-14 09:45:12 0 [IP Address]    100
## 8 2019-06-14 09:43:27 2019-06-14 09:45:12 0 [IP Address]    100
## 9 2019-06-14 09:43:08 2019-06-14 09:45:13 0 [IP Address]    100
## 10 2019-06-14 09:43:36 2019-06-14 09:45:16 0 [IP Address]    100
```

```
## # ... with 1,450 more rows, and 46 more variables:
```

```
## #   Duration_in_seconds_ <dbl>, Finished <dbl+lbl>,
## #   RecordedDate <dtm>, ResponseId <chr>, UserLanguage <chr>,
## #   Pol_7 <dbl+lbl>, W2_Knowledge <dbl+lbl>, Gender <dbl+lbl>,
## #   Race <dbl+lbl>, W1_Feeling_1 <dbl>, W1_Actions_1_1 <dbl+lbl>,
## #   W1_Actions_1_2 <dbl+lbl>, W1_Actions_1_3 <dbl+lbl>,
```


The pipe

So far, we have written our code like this:

```
filter(ads_data, Duration__in_seconds_ < 100)
```

```
## # A tibble: 41 × 52
##   StartDate      EndDate      Status Progress
##   <dtm>         <dtm>         <dbl+lbl>    <dbl>
## 1 2019-06-14 09:43:20 2019-06-14 09:44:30 0 [IP Address]    100
## 2 2019-06-14 09:43:29 2019-06-14 09:44:58 0 [IP Address]    100
## 3 2019-06-14 09:44:00 2019-06-14 09:45:11 0 [IP Address]    100
## 4 2019-06-14 09:43:32 2019-06-14 09:45:12 0 [IP Address]    100
## 5 2019-06-14 09:43:48 2019-06-14 09:45:25 0 [IP Address]    100
## 6 2019-06-14 09:44:24 2019-06-14 09:45:26 0 [IP Address]    100
## 7 2019-06-14 09:43:50 2019-06-14 09:45:29 0 [IP Address]    100
## 8 2019-06-14 09:44:15 2019-06-14 09:45:42 0 [IP Address]    100
## 9 2019-06-14 09:44:30 2019-06-14 09:45:58 0 [IP Address]    100
## 10 2019-06-14 09:44:36 2019-06-14 09:46:05 0 [IP Address]    100
## # ... with 31 more rows, and 48 more variables:
## #   Duration__in_seconds_ <dbl>, Finished <dbl+lbl>,
## #   RecordedDate <dtm>, ResponseId <chr>, DistributionChannel <chr>,
## #   UserLanguage <chr>, Consent <dbl+lbl>, Pol_7 <dbl+lbl>,
## #   W2_Knowledge <dbl+lbl>, Gender <dbl+lbl>, Race <dbl+lbl>,
## #   W1_Feeling_1 <dbl>, W1_Actions_1_1 <dbl+lbl>,
```

The pipe

We can use the pipe `%>%`, which passes what we wrote on the previous line into the next function as the first argument:

```
ads_data %>%  
  filter(Duration__in_seconds_ < 100) %>%  
  arrange(Duration__in_seconds_) %>%  
  select(RecordedDate, Duration__in_seconds_)
```

```
## # A tibble: 41 × 2  
##   RecordedDate      Duration__in_seconds_  
##   <dtm>              <dbl>  
## 1 2019-06-14 09:59:02          50  
## 2 2019-06-14 09:45:26          61  
## 3 2019-06-14 09:44:31          70  
## 4 2019-06-14 09:45:12          70  
## 5 2019-06-14 09:53:26          75  
## 6 2019-06-14 09:47:13          76  
## 7 2019-06-14 09:51:54          76  
## 8 2019-06-14 09:47:08          78  
## 9 2019-06-14 10:11:46          79  
## 10 2019-06-14 09:54:54          80  
## # ... with 31 more rows
```

The pipe

```
ads_data %>%  
  filter(Duration__in_seconds_ < 100) %>%  
  arrange(Duration__in_seconds_) %>%  
  select(RecordedDate, Duration__in_seconds_)
```

You can think of this like:

- Take the ADS data
- Filter so we only have the rows where the survey duration is less than 100 seconds
- Arrange so we go from lowest duration to highest
- Select only the date recorded and the duration

Mutating

Mutating can be used to create new columns or change existing columns.

```
ads_data <- ads_data %>%  
  mutate(Birthyear_add_day = str_c(Birthyear, "07-01")) %>%  
  mutate(Birthyear_add_day = as_datetime(Birthyear_add_day))
```

```
## # A tibble: 1,460 × 3  
##   EndDate          Birthyear Birthyear_add_day  
##   <dtm>          <dbl> <dtm>  
## 1 2019-06-14 09:44:30      1993 1993-07-01 00:00:00  
## 2 2019-06-14 09:44:57      1978 1978-07-01 00:00:00  
## 3 2019-06-14 09:44:58      1993 1993-07-01 00:00:00  
## 4 2019-06-14 09:45:00      1983 1983-07-01 00:00:00  
## 5 2019-06-14 09:45:00      1990 1990-07-01 00:00:00  
## 6 2019-06-14 09:45:11      1980 1980-07-01 00:00:00  
## 7 2019-06-14 09:45:12      1996 1996-07-01 00:00:00  
## 8 2019-06-14 09:45:12      1986 1986-07-01 00:00:00  
## 9 2019-06-14 09:45:13      2000 2000-07-01 00:00:00  
## 10 2019-06-14 09:45:16      1988 1988-07-01 00:00:00  
## # ... with 1,450 more rows
```

Mutating

```
ads_data %>%  
  mutate(age = EndDate - Birthyear_add_day)
```

```
## # A tibble: 1,460 × 4  
##   EndDate      Birthyear Birthyear_add_day    age  
##   <dtm>          <dbl> <dtm>          <drtn>  
## 1 2019-06-14 09:44:30    1993 1993-07-01 00:00:00  9479.406 days  
## 2 2019-06-14 09:44:57    1978 1978-07-01 00:00:00 14958.406 days  
## 3 2019-06-14 09:44:58    1993 1993-07-01 00:00:00  9479.406 days  
## 4 2019-06-14 09:45:00    1983 1983-07-01 00:00:00 13132.406 days  
## 5 2019-06-14 09:45:00    1990 1990-07-01 00:00:00 10575.406 days  
## 6 2019-06-14 09:45:11    1980 1980-07-01 00:00:00 14227.406 days  
## 7 2019-06-14 09:45:12    1996 1996-07-01 00:00:00  8383.406 days  
## 8 2019-06-14 09:45:12    1986 1986-07-01 00:00:00 12036.406 days  
## 9 2019-06-14 09:45:13    2000 2000-07-01 00:00:00  6922.406 days  
## 10 2019-06-14 09:45:16    1988 1988-07-01 00:00:00 11305.406 days  
## # ... with 1,450 more rows
```

Summary

```
summary(ads_data)
```

```
##      StartDate                      EndDate
##  Min.      :2019-06-14 09:43:03.00  Min.      :2019-06-14 09:44:30.00
##  1st Qu.:2019-06-14 09:46:47.50  1st Qu.:2019-06-14 09:51:29.00
##  Median :2019-06-14 09:52:50.00  Median :2019-06-14 09:57:57.00
##  Mean    :2019-06-14 09:57:40.11  Mean    :2019-06-14 10:02:23.89
##  3rd Qu.:2019-06-14 10:06:28.25  3rd Qu.:2019-06-14 10:11:19.50
##  Max.    :2019-06-14 11:19:45.00  Max.    :2019-06-14 11:27:10.00
##
##      Status      Progress      Duration__in_seconds_      Finished
##  Min.      :0      Min.      :100      Min.      : 50.0      Min.      :1
##  1st Qu.:0      1st Qu.:100      1st Qu.: 178.0      1st Qu.:1
##  Median :0      Median :100      Median : 237.0      Median :1
##  Mean    :0      Mean    :100      Mean    : 283.3      Mean    :1
##  3rd Qu.:0      3rd Qu.:100      3rd Qu.: 324.2      3rd Qu.:1
##  Max.    :0      Max.    :100      Max.    :1575.0      Max.    :1
##
##      RecordedDate                      ResponseId
##  Min.      :2019-06-14 09:44:31.00  Length:1460
##  1st Qu.:2019-06-14 09:51:29.00  Class :character
##  Median :2019-06-14 09:57:58.00  Mode  :character
##  Mean    :2019-06-14 10:02:24.49
```

Pulling a variable for calculations

```
ads_data %>%  
  pull(Duration__in_seconds_)
```

```
##      [1]      70      105      88      109      109      70      99      105      124      100      96      102      61  
##     [14]      98      120      86      119      120      143      115      131      164      140      126      88      127  
##     [27]     146       88     134     163     111     164     123     176     102     119     187     179     140  
##     [40]     144     183     139     123     162     152     184     160     181     163     168     101     190  
##     [53]     178     144     194     123     133     135     185     121     163     192     210     167     139  
##     [66]     204     117     170     170     199       95     126     208     178     207     146     118     170  
##     [79]     110     172     226       78     160     185     186     222     212     185     168     213       76  
##     [92]     213     165     173     218     207     214     203     206     213     228     186     240     248  
##    [105]     208     176     217     142     190     215     247     163     239     251     185     176     217  
##    [118]     193     171     159     239     252     178     168     101     213     227     122     217     225  
##    [131]     239     182     178     165     248     190     272     222     101     173     270     121     191  
##    [144]     275     210     227     283     188     194     275     236     169     151     295     262     257  
##    [157]     234     119     287     276     264     286     193     245     196     289     148     295     208  
##    [170]     285     209     318     210     113     193     262     322     168     298     278     216     228  
##    [183]     252     185     343     121     319     281     239     115     321     303     304     300     267  
##    [196]     190     228     194     271     187     283     232     164     241     213     288     188     323  
##    [209]     237     265     245     174     361     172     276     195     357     226     188     223     234  
##    [222]     291     197     283     339     100     319     216     224     169     182     257     227     347
```

Using the pulled variable for descriptive statistics

Median

```
ads_data %>%  
  pull(Duration__in_seconds_) %>%  
  median(na.rm = TRUE)
```

```
## [1] 237
```

We have to tell the `mean()` function to disregard NAs by writing `na.rm = TRUE`

Using the pulled variable for descriptive statistics

Mean

```
ads_data %>%  
  pull(Duration__in_seconds_) %>%  
  mean(na.rm = TRUE)
```

```
## [1] 283.261
```

Using the pulled variable for descriptive statistics

Range can be calculated using the `range()` function.

```
ads_data %>%  
  pull(Duration__in_seconds_) %>%  
  range(na.rm = TRUE)
```

```
## [1]    50 1575
```

Variance can be calculated using the `var()` function.

```
ads_data %>%  
  pull(Duration__in_seconds_) %>%  
  var(na.rm = TRUE)
```

```
## [1] 29487.81
```

Using the pulled variable for descriptive statistics

Standard Deviation can be calculated using the `sd()` function.

```
ads_data %>%  
  pull(Duration__in_seconds_) %>%  
  sd(na.rm = TRUE)
```

```
## [1] 171.7202
```

Summarise

```
ads_data %>%  
  summarise(mean_time = mean(Duration__in_seconds_, na.rm = TRUE),  
            sd_time = sd(Duration__in_seconds_, na.rm = TRUE))
```

```
## # A tibble: 1 × 2  
##   mean_time sd_time  
##   <dbl>    <dbl>  
## 1     283.     172.
```

Grouping

Before summarising, we can group by a categorical variable

```
ads_data %>%  
  group_by(Gender) %>%  
  summarise(count = n(),  
            mean_time = mean(Duration__in_seconds_, na.rm = TRUE),  
            sd_time = sd(Duration__in_seconds_, na.rm = TRUE))
```

```
## # A tibble: 3 × 4
```

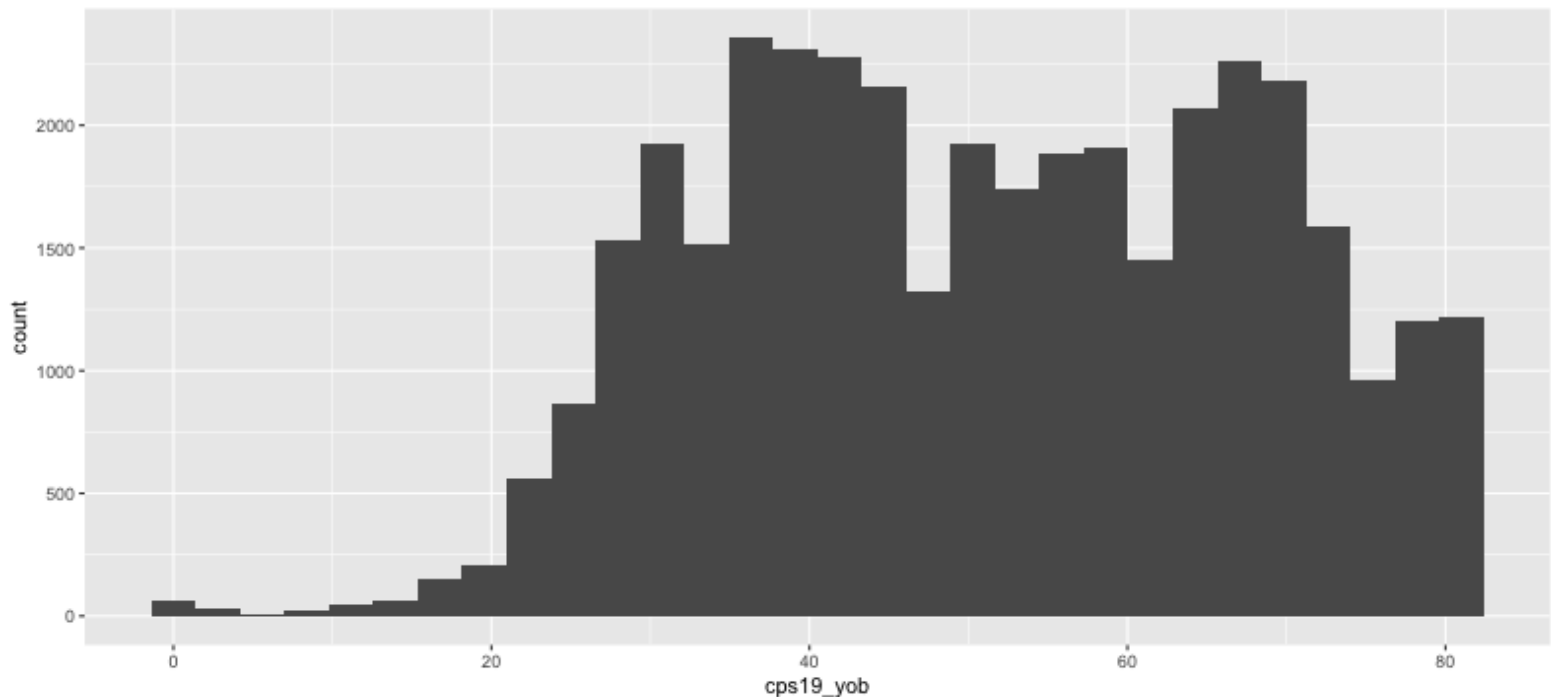
```
##           Gender count mean_time sd_time  
##      <dbl+lbl> <int>    <dbl>    <dbl>  
## 1 1 [Male]      758      269.     162.  
## 2 2 [Female]    698      299.     181.  
## 3 3 [Prefer a third option/Other] 4      229     37.7
```

Manipulation application: data cleaning

Data cleaning

Graphing year of birth shows that it goes from 1 to about 80.

```
ces_2019_raw %>%  
  ggplot(aes(x = cps19_yob)) +  
  geom_histogram()
```



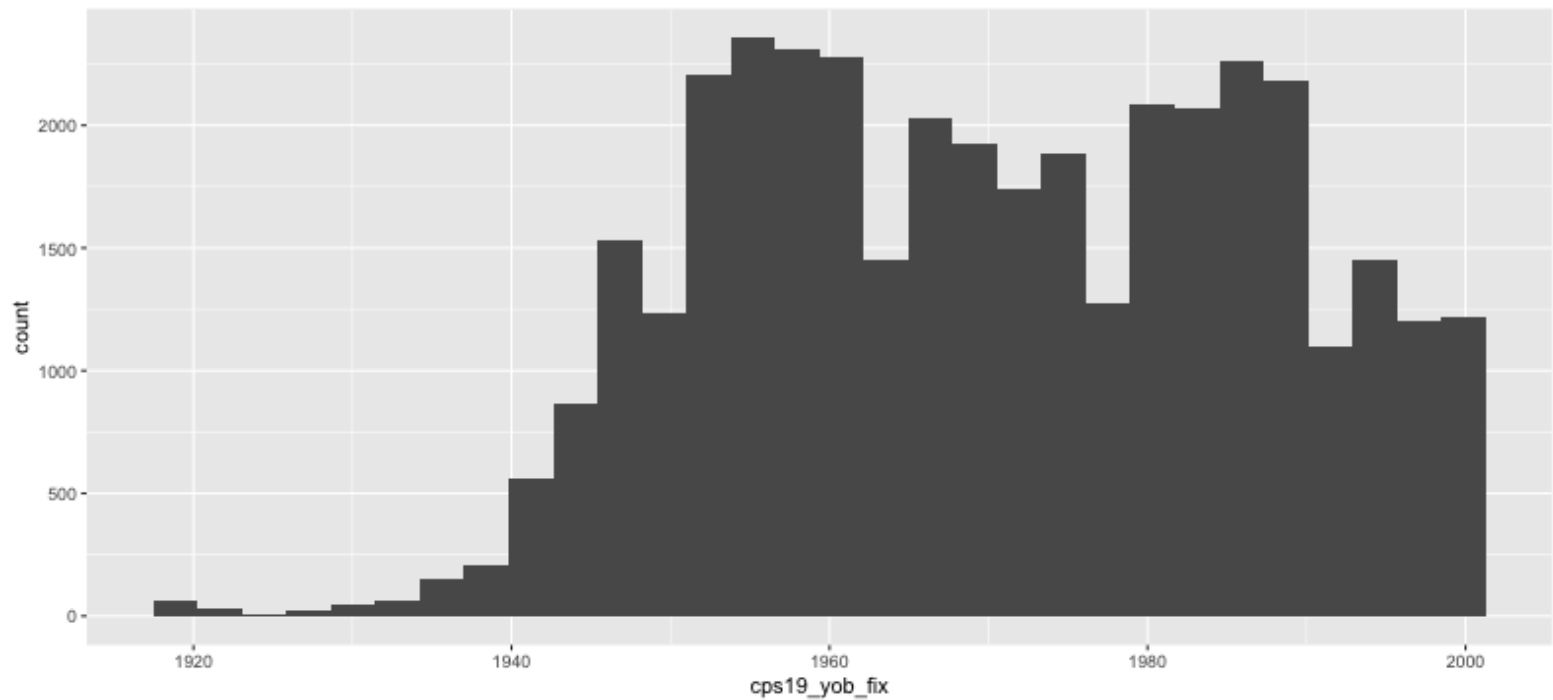
Data cleaning

The codebook says that a value of 1 corresponds to a birth year of 1920, value of 2 to a birth year of 1921, and so on. We can create a new variable that reads more intuitively.

```
CES_data <- ces_2019_raw %>%  
  mutate(cps19_yob_fix = cps19_yob + 1919)
```


Data cleaning

```
CES_data %>%  
  ggplot(aes(x = cps19_yob_fix)) +  
  geom_histogram()
```



Better!

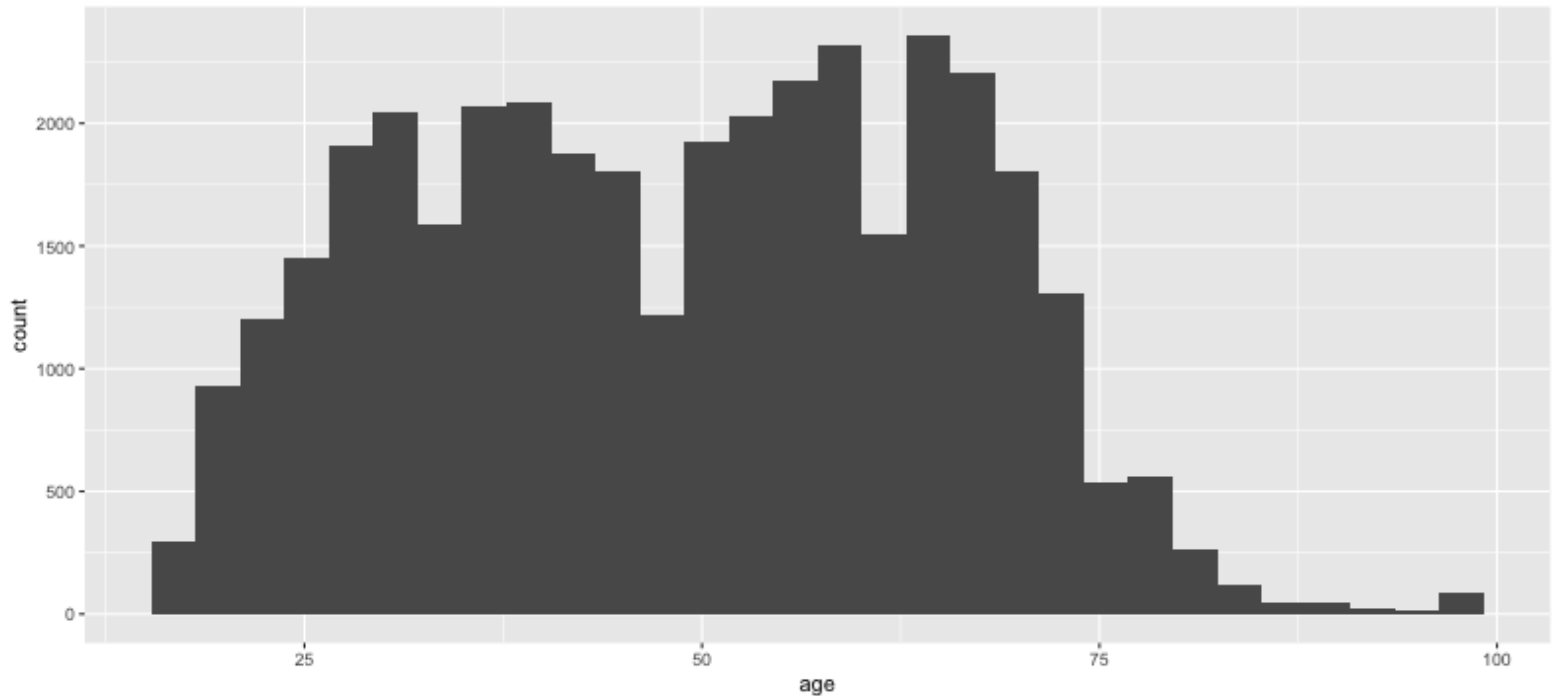
Add a variable for age

Now that we have an accurate birth year, maybe we would like to have the age of the individual as well.

```
CES_data <- CES_data %>%  
  mutate(age = 2019 - cps19_yob_fix)
```

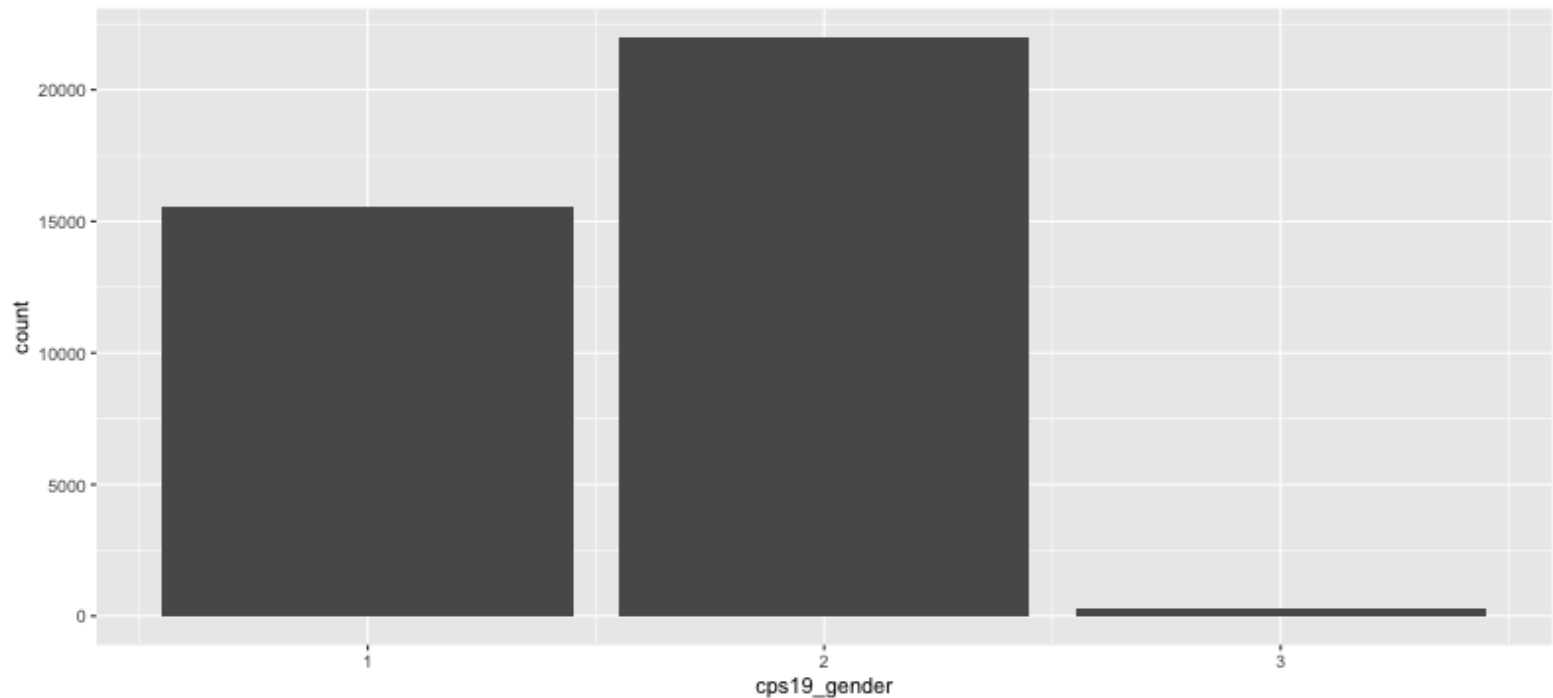
Add a variable for age

```
CES_data %>%  
  ggplot(aes(x = age)) +  
  geom_histogram()
```



Recoding the gender variable

```
CES_data %>%  
  ggplot(aes(x = cps19_gender)) +  
  geom_bar()
```




Recoding the gender variable

```
CES_data <- CES_data %>%  
  mutate(cps19_gender_fix = factor(cps19_gender)) %>%  
  mutate(cps19_gender_fix = fct_recode(cps19_gender_fix,  
                                       "M" = "1",  
                                       "F" = "2",  
                                       "NB" = "3"))
```

Recoding the gender variable

```
CES_data %>%  
  ggplot(aes(x = cps19_gender_fix)) +  
  geom_bar()
```



Fixing household counts

```
CES_data %>%  
  filter(cps19_household > 10) %>%  
  arrange(-cps19_household) %>%  
  pull(cps19_household)
```

##	[1]	7766666	72000	50000	20000	10000	5667	2000	501
##	[9]	321	99	89	87	69	54	54	50
##	[17]	44	40	34	33	29	27	23	22
##	[25]	22	20	20	20	15	15	13	13
##	[33]	12	12	12	11	11	11	11	11
##	[41]	11	11	11					

Fixing household counts

```
CES_data <- CES_data %>%  
  mutate(cps19_household = ifelse(cps19_household > 15,  
                                   NA,  
                                   cps19_household))
```

```
CES_data %>%  
  filter(cps19_household > 10) %>%  
  pull(cps19_household)
```

```
## [1] 12 11 15 12 11 13 11 11 11 15 13 12 11 11 11
```


Fixing income

```
CES_data %>%  
  filter(cps19_income_number > 1000000) %>%  
  arrange(-cps19_income_number) %>%  
  pull(cps19_income_number)
```

```
## [1] 6.747658e+60 1.000000e+21 1.000000e+15 8.769655e+10 8.889899e+09  
## [6] 3.062936e+09 1.000000e+09 1.000000e+09 6.788765e+08 3.000000e+08  
## [11] 7.245600e+07 3.454534e+07 3.000000e+07 1.000000e+07 9.999999e+06  
## [16] 8.900000e+06 7.696588e+06 7.440000e+06 6.848382e+06 6.787145e+06  
## [21] 6.782800e+06 6.500100e+06 4.500000e+06 3.000000e+06 2.332100e+06  
## [26] 2.000000e+06 2.000000e+06 1.872717e+06 1.800000e+06 1.650000e+06  
## [31] 1.500000e+06 1.500000e+06 1.450000e+06 1.300000e+06 1.290000e+06  
## [36] 1.250000e+06 1.250000e+06 1.250000e+06 1.150000e+06
```

Fixing income

```
CES_data <- CES_data %>%  
  mutate(cps19_income_number = ifelse(cps19_income_number >= 1000000000,  
                                       NA,  
                                       cps19_income_number))  
  
CES_data %>%  
  filter(cps19_income_number > 1000000) %>%  
  pull(cps19_income_number)
```

```
## [1] 2000000 1500000 4500000 3000000 6848382 7696588  
## [7] 6787145 1250000 1650000 1872717 678876545 1300000  
## [13] 1150000 1250000 9999999 1450000 1500000 6500100  
## [19] 300000000 8900000 300000000 7440000 6782800 2332100  
## [25] 1800000 2000000 10000000 1290000 72456000 34545345  
## [31] 1250000
```

Manipulation application: Summarising data

Summarising data

First we can select only data for Ontario using `filter()`:

```
CES_data %>%  
  filter(cps19_province == "Ontario")
```

```
## # A tibble: 14,160 × 620  
##   cps19_StartDate      cps19_EndDate      cps19_ResponseId  
##   <dtm>              <dtm>              <chr>  
## 1 2019-09-13 10:01:19 2019-09-13 10:27:29 R_USWDAPcQEQiMmNb  
## 2 2019-09-13 10:05:37 2019-09-13 10:50:53 R_3IQaeDXy0tBzEry  
## 3 2019-09-13 10:05:52 2019-09-13 10:32:53 R_27WeMQ1asip2cMD  
## 4 2019-09-13 10:10:20 2019-09-13 10:29:45 R_3LiGZcCWJecWV4P  
## 5 2019-09-13 10:14:47 2019-09-13 10:32:32 R_1Iu8R1UlyzVMycz  
## 6 2019-09-13 10:15:39 2019-09-13 10:30:59 R_2EcS26hqrcVYlab  
## 7 2019-09-13 10:15:48 2019-09-13 10:37:45 R_3yrt44wqQ1d4VRn  
## 8 2019-09-13 10:16:08 2019-09-13 10:40:14 R_100BmXJyvn8feYQ  
## 9 2019-09-13 10:16:24 2019-09-13 10:41:24 R_2e5nvu0UchQctgq  
## 10 2019-09-13 10:17:06 2019-09-13 10:35:47 R_20Jdv16hkRGnjOn  
## # ... with 14,150 more rows, and 617 more variables:  
## #   cps19_consent <dbl>, cps19_citizenship <dbl>, cps19_yob <dbl>,  
## #   cps19_yob_2001_age <dbl>, cps19_gender <fct>,  
## #   cps19_province <fct>, cps19_education <dbl>, cps19_demsat <dbl>,  
## #   cps19_imp_iss <chr>, cps19_imp_iss_party <dbl>,
```

Summarising data

We don't need to be dealing with all the columns. We can specifically select the ones we want using `select()`:

"How satisfied are you with the performance of your provincial government under \${e://Field/premier}?", "In provincial politics, do you usually think of yourself as a:", and income.

```
CES_data %>%  
  filter(cps19_province == "Ontario") %>%  
  select(cps19_prov_gov_sat,  
         cps19_prov_id,  
         cps19_income_number)
```

```
## # A tibble: 14,160 × 3  
##   cps19_prov_gov_sat cps19_prov_id cps19_income_number  
##   <fct>             <fct>             <dbl>  
## 1 Not very satisfied Liberal                NA  
## 2 Fairly satisfied  Progressive Conservative NA  
## 3 Fairly satisfied  Liberal                56000  
## 4 Not at all satisfied NDP                NA  
## 5 Not at all satisfied NDP                0  
## 6 Not at all satisfied None                NA  
## 7 Not at all satisfied NDP                NA
```

Summarising data

Now that our data looks like what we would like it to, we can start creating a summary table. Since we have the income for each participant, we can look at median incomes. We also want to know how many participants are in each category.

First, we can group the data by provincial political self-ID. To do this, we use `group_by()` to group the data and `summarise()` to produce values for each group we have created. We will start with calculating the `median()` for the incomes. We can add multiple arguments to the `summarise()` argument. `n()` adds a count for each group.

Summarising data

```
CES_data %>%  
  filter(cps19_province == "Ontario") %>%  
  select(cps19_prov_gov_sat,  
         cps19_prov_id,  
         cps19_income_number) %>%  
  group_by(cps19_prov_gov_sat) %>%  
  summarise(median_income = median(cps19_income_number,  
                                   na.rm = TRUE),  
            count = n())
```

```
## # A tibble: 5 × 3  
##   cps19_prov_gov_sat median_income count  
##   <fct>              <dbl> <int>  
## 1 Very satisfied    80000    872  
## 2 Fairly satisfied  80000   2738  
## 3 Not very satisfied 75000   3212  
## 4 Not at all satisfied 72000   6853  
## 5 Don't know/prefer not to answer 50000    485
```

Grouping

In our table, the satisfaction ratings are ordered alphabetically. We would like them to be ordered logically. We can do this by ordering the factor variable.

```
CES_data %>%  
  filter(cps19_province == "Ontario") %>%  
  select(cps19_prov_gov_sat,  
         cps19_prov_id,  
         cps19_income_number) %>%  
  mutate(cps19_prov_id = factor(cps19_prov_id,  
                                levels = c("Liberal",  
                                             "Progressive Conservative",  
                                             "NDP",  
                                             "Green",  
                                             "Another party",  
                                             "None",  
                                             "Don't know/prefer not to say")))
```


Grouping

```
## # A tibble: 14,160 × 3
##   cps19_prov_gov_sat  cps19_prov_id  cps19_income_number
##   <fct>              <fct>              <dbl>
## 1 Not very satisfied Liberal                NA
## 2 Fairly satisfied   Progressive Conservative NA
## 3 Fairly satisfied   Liberal                56000
## 4 Not at all satisfied NDP                NA
## 5 Not at all satisfied NDP                 0
## 6 Not at all satisfied None                NA
## 7 Not at all satisfied NDP                NA
## 8 Not very satisfied Liberal                NA
## 9 Not very satisfied NDP                NA
## 10 Not at all satisfied Liberal            NA
## # ... with 14,150 more rows
```

Grouping

And combine this with our table from before:

```
CES_data %>%
  filter(cps19_province == "Ontario") %>%
  select(cps19_prov_gov_sat,
         cps19_prov_id,
         cps19_income_number) %>%
  mutate(cps19_prov_gov_sat = factor(cps19_prov_gov_sat,
                                     levels = c("Not at all satisfied",
                                                "Not very satisfied",
                                                "Fairly satisfied",
                                                "Very satisfied",
                                                "Don't know/prefer not to answer"),
                                     ordered = TRUE))

group_by(cps19_prov_gov_sat) %>%
  summarise(median_income = median(cps19_income_number,
                                   na.rm = TRUE),
            count = n())
```

```
## # A tibble: 5 × 3
##   cps19_prov_gov_sat median_income count
##   <fct>              <dbl> <int>
## 1 Not at all satisfied    72000   6853
## 2 Not very satisfied     75000   3212
```

Grouping

What happens if we group by political identification instead?

```
CES_data %>%  
  filter(cps19_province == "Ontario") %>%  
  select(cps19_prov_gov_sat,  
         cps19_prov_id,  
         cps19_income_number) %>%  
  group_by(cps19_prov_id) %>%  
  summarise(median_income = median(cps19_income_number,  
                                   na.rm = TRUE),  
            count = n())
```

```
## # A tibble: 7 × 3  
##   cps19_prov_id median_income count  
##   <fct>          <dbl> <int>  
## 1 Liberal      80000  4607  
## 2 NDP          65000  2413  
## 3 Green        60000   812  
## 4 Progressive Conservative 80000  3629  
## 5 Another party 50000   90  
## 6 None         68000 1367  
## 7 Don't know/prefer not to answer 60000 1242
```

Grouping

We could order the parties in a way that makes more sense:

```
CES_data %>%
  filter(cps19_province == "Ontario") %>%
  select(cps19_prov_gov_sat,
         cps19_prov_id,
         cps19_income_number) %>%
  mutate(cps19_prov_id = factor(cps19_prov_id,
                                levels = c("Liberal",
                                             "Progressive Conservative",
                                             "NDP",
                                             "Green",
                                             "Another party",
                                             "None",
                                             "Don't know/prefer not to say"))

group_by(cps19_prov_id) %>%
  summarise(median_income = median(cps19_income_number,
                                   na.rm = TRUE),
            count = n())
```

Grouping

```
## # A tibble: 7 × 3
##   cps19_prov_id median_income count
##   <fct>          <dbl> <int>
## 1 Liberal        80000   4607
## 2 Progressive    80000   3629
## 3 NDP            65000   2413
## 4 Green          60000    812
## 5 Another party  50000     90
## 6 None           68000   1367
## 7 Don't know/prefer not to answer 60000   1242
```

Grouping

Or we could sort by median income. We can do that using `arrange()`:

```
CES_data %>%
  filter(cps19_province == "Ontario") %>%
  select(cps19_prov_gov_sat,
         cps19_prov_id,
         cps19_income_number) %>%
  group_by(cps19_prov_id) %>%
  summarise(median_income = median(cps19_income_number,
                                   na.rm = TRUE),
            count = n()) %>%
  arrange(-median_income)
```

```
## # A tibble: 7 × 3
##   cps19_prov_id median_income count
##   <fct>          <dbl> <int>
## 1 Liberal          80000    4607
## 2 Progressive Conservative 80000    3629
## 3 None             68000    1367
## 4 NDP             65000    2413
## 5 Green           60000     812
## 6 Don't know/prefer not to answer 60000    1242
## 7 Another party    50000     90
```

Grouping

`group_by()` can also have multiple arguments, so we can group by `cps19_prov_gov_sat` and `cps19_prov_id` at the same time:

```
CES_data %>%
  filter(cps19_province == "Ontario") %>%
  select(cps19_prov_gov_sat,
         cps19_prov_id,
         cps19_income_number) %>%
  mutate(cps19_prov_id = factor(cps19_prov_id,
                                levels = c("Liberal",
                                             "Progressive Conservative",
                                             "NDP",
                                             "Green",
                                             "Another party",
                                             "None",
                                             "Don't know/prefer not to answer"),
         cps19_prov_gov_sat = factor(cps19_prov_gov_sat,
                                      levels = c("Not at all satisfied",
                                                  "Not very satisfied",
                                                  "Fairly satisfied",
                                                  "Very satisfied",
                                                  "Don't know/prefer not to answer"),
  group_by(cps19_prov_gov_sat, cps19_prov_id) %>%
```

Grouping

This table is less easy to read, though. `spread()` can make a table that is wide rather than long. We specify the **key**, the variable that will become our column names, and the **value**, which will become the values in those columns:

```
CES_data %>%
  filter(cps19_province == "Ontario") %>%
  select(cps19_prov_gov_sat,
         cps19_prov_id,
         cps19_income_number) %>%
  mutate(cps19_prov_id = factor(cps19_prov_id,
                                levels = c("Liberal",
                                             "Progressive Conservative",
                                             "NDP",
                                             "Green",
                                             "Another party",
                                             "None",
                                             "Don't know/prefer not to answer"),
         cps19_prov_gov_sat = factor(cps19_prov_gov_sat,
                                      levels = c("Not at all satisfied",
                                                  "Not very satisfied",
                                                  "Fairly satisfied",
                                                  "Very satisfied",
                                                  "Don't know/prefer not to answer"))
```


Grouping

```
## # A tibble: 7 × 6
##   cps19_prov_id      `Not at all sa...` `Not very sati...` `Fairly satisf...`
##   <fct>                <dbl>                <dbl>                <dbl>
## 1 Liberal              80000              80000              79999
## 2 Progressive Conse...  85000              78000              82000
## 3 NDP                  65000              65000              76888
## 4 Green                 60000              60000              72750
## 5 Another party        40000              48500              73500
## 6 None                  62000              74000              69000
## 7 Don't know/prefer...  68500              59500              70000
## # ... with 2 more variables: `Very satisfied` <dbl>,
## #   `Don't know/prefer not to answer` <dbl>
```

Exercises

Exercises

1. Filter the rows in the CES_data dataset where the survey-taker is between 30 and 50 (cps19_age).
2. Filter the rows in the CES_data dataset where the survey-taker answered the cps19_votchoice question (i.e. the cps19_votchoice variable is not NA).
3. Select the variables cps19_age and cps19_province from the CES_data dataset.
4. Select all variables except cps19_province from the CES_data dataset.

Exercises

1. Create a variable in the dataset CES_data that states if a person consumes news content or not (i.e. cps19_news_cons is equal to "0 minutes" or it is not).
2. Modify the variable cps19_income_number in the dataset CES_data so that it is measured in thousands (i.e. divide the income number by 1000).

Exercises

1. Use the CES_data dataset. Group by cps19_votchoice. Find both the median and mean rating of Trudeau (cps19_lead_rating_23):
2. Use the CES_data dataset. Group by cps19_imm and cps19_spend_educ. Find the count for each group.

Exercises

- 1 - Fix this error:

```
CES_data %>%  
  summarise(mean = mean(cps19_age)) %>%  
  group_by(cps19_gender)
```

- 2 - Fix this error:

```
CES_data %>%  
  filter(cps19_vote_choice == "Green Party")
```

Exercises

- 3 - Fix this error:

```
CES_data %>%  
  mutate(cps19_fed_donate = factor(cps19_fed_donate,  
                                   levels = c("Yes",  
                                              "No",  
                                              "Don't know/ Prefer not
```

- 4 - Fix this error:

```
CES_data %>%  
  select(cps19_province  
         cps19_age  
         cps19_gender)
```

Any questions?